

Math 362: Mathematical Statistics II

Le Chen

le.chen@emory.edu
chenle02@gmail.com

Emory University
Atlanta, GA

Last updated on Spring 2021
Last compiled on January 15, 2023

2021 Spring

Creative Commons License
(CC By-NC-SA)

Chapter 5. Estimation

§ 5.1 Introduction

§ 5.2 Estimating parameters: MLE and MME

§ 5.3 Interval Estimation

§ 5.4 Properties of Estimators

§ 5.5 Minimum-Variance Estimators: The Cramér-Rao Lower Bound

§ 5.6 Sufficient Estimators

§ 5.7 Consistency

§ 5.8 Bayesian Estimation

Chapter 5. Estimation

§ 5.1 Introduction

§ 5.2 Estimating parameters: MLE and MME

§ 5.3 Interval Estimation

§ 5.4 Properties of Estimators

§ 5.5 Minimum-Variance Estimators: The Cramér-Rao Lower Bound

§ 5.6 Sufficient Estimators

§ 5.7 Consistency

§ 5.8 Bayesian Estimation

Motivating example: Given an unfair coin, or p -coin, such that

$$X = \begin{cases} 1 & \text{head with probability } p, \\ 0 & \text{tail with probability } 1 - p, \end{cases}$$

how would you determine the value p ?


Solutions:

1. You need to try the coin several times, say, three times. What you obtain is “HHT”.
2. Draw a conclusion from the experiment you just made.

Rationale: The choice of the parameter p should be the value that maximizes the probability of the sample.

$$\begin{aligned}\mathbb{P}(X_1 = 1, X_2 = 1, X_3 = 0) &= P(X_1 = 1)P(X_2 = 1)P(X_3 = 0) \\ &= p^2(1 - p).\end{aligned}$$

```
1 # Hello, R.  
2 p <- seq(0,1,0.01)  
3 plot(p,p^2*(1-p),  
4     type="l",  
5     col="red")  
6 title("Likelihood")  
7 # add a vertical dotted (4) blue line  
8 abline(v=0.67, col="blue", lty=4)  
9 # add some text  
10 text(0.67,0.01, "2/3")
```



Maximize $f(p) = p^2(1 - p) \dots$

A random sample of size n from the population – Bernoulli(p):

- ▶ X_1, \dots, X_n are i.i.d.¹ random variables, each following Bernoulli(p).
- ▶ Suppose the outcomes of the random sample are: $X_1 = k_1, \dots, X_n = k_n$.
- ▶ What is your choice of p based on the above random sample?

$$\hat{p} = \frac{1}{n} \sum_{i=1}^n k_i =: \bar{k}.$$

¹independent and identically distributed

A random sample of size n from the population with given pdf:

- ▶ X_1, \dots, X_n are i.i.d. random variables, each following the same given pdf.
- ▶ a **statistic** or an **estimator** is a function of the random sample.

Statistic/Estimator is a random variable!

e.g.,

$$\hat{p} = \frac{1}{n} \sum_{i=1}^n X_i.$$

- ▶ The outcome of a statistic/estimator is called an **estimate**. e.g.,

$$p_e = \frac{1}{n} \sum_{i=1}^n k_i.$$

Chapter 5. Estimation

§ 5.1 Introduction

§ 5.2 Estimating parameters: MLE and MME

§ 5.3 Interval Estimation

§ 5.4 Properties of Estimators

§ 5.5 Minimum-Variance Estimators: The Cramér-Rao Lower Bound

§ 5.6 Sufficient Estimators

§ 5.7 Consistency

§ 5.8 Bayesian Estimation

Two methods for estimating parameters

Corresponding estimator

1. Method of maximum likelihood.

MLE

2. Method of moments.

MME

Maximum Likelihood Estimation

Definition 5.2.1. For a random sample of size n from the discrete (resp. continuous) population/pdf $p_X(k; \theta)$ (resp. $f_Y(y; \theta)$), the **likelihood function**, $L(\theta)$, is the product of the pdf evaluated at $X_i = k_i$ (resp. $Y_i = y_i$), i.e.,

$$L(\theta) = \prod_{i=1}^n p_X(k_i; \theta) \quad \left(\text{resp. } L(\theta) = \prod_{i=1}^n f_Y(y_i; \theta) \right).$$

Definition 5.2.2. Let $L(\theta)$ be as defined in Definition 5.2.1. If θ_e is a value of the parameter such that $L(\theta_e) \geq L(\theta)$ for all possible values of θ , then we call θ_e the **maximum likelihood estimate** for θ .

Examples for MLE

Often but not always MLE can be obtained by setting the first derivative equal to zero:

E.g. 1. Poisson distribution: $p_X(k) = e^{-\lambda} \frac{\lambda^k}{k!}$, $k = 0, 1, \dots$.

$$L(\lambda) = \prod_{i=1}^n e^{-\lambda} \frac{\lambda^{k_i}}{k_i!} = e^{-n\lambda} \lambda^{\sum_{i=1}^n k_i} \left(\prod_{i=1}^n k_i! \right)^{-1}.$$

$$\ln L(\lambda) = -n\lambda + \left(\sum_{i=1}^n k_i \right) \ln \lambda - \ln \left(\prod_{i=1}^n k_i! \right).$$

$$\frac{d}{d\lambda} \ln L(\lambda) = -n + \frac{1}{\lambda} \sum_{i=1}^n k_i.$$

$$\frac{d}{d\lambda} \ln L(\lambda) = 0 \quad \Rightarrow \quad \boxed{\lambda_e = \frac{1}{n} \sum_{i=1}^n k_i =: \bar{k}}.$$

Comment: The critical point is indeed global maximum because

$$\frac{d^2}{d\lambda^2} \ln L(\lambda) = -\frac{1}{\lambda^2} \sum_{i=1}^n k_i < 0.$$

The following two cases are related to waiting time:

E.g. 2. Exponential distribution: $f_Y(y) = \lambda e^{-\lambda y}$ for $y \geq 0$.

$$L(\lambda) = \prod_{i=1}^n \lambda e^{-\lambda y_i} = \lambda^n \exp \left(-\lambda \sum_{i=1}^n y_i \right)$$

$$\ln L(\lambda) = n \ln \lambda - \lambda \sum_{i=1}^n y_i.$$

$$\frac{d}{d\lambda} \ln L(\lambda) = \frac{n}{\lambda} - \sum_{i=1}^n y_i.$$

$$\frac{d}{d\lambda} \ln L(\lambda) = 0 \quad \Longrightarrow \quad \boxed{\lambda_e = \frac{n}{\sum_{i=1}^n y_i} =: \frac{1}{\bar{y}}}.$$

A random sample of size n from the following population:

E.g. 3. Gamma distribution: $f_Y(y; \lambda) = \frac{\lambda^r}{\Gamma(r)} y^{r-1} e^{-\lambda y}$ for $y \geq 0$ with $r > 1$ known.

$$L(\lambda) = \prod_{i=1}^n \frac{\lambda^r}{\Gamma(r)} y_i^{r-1} e^{-\lambda y_i} = \lambda^{r n} \Gamma(r)^{-n} \left(\prod_{i=1}^n y_i^{r-1} \right) \exp \left(-\lambda \sum_{i=1}^n y_i \right)$$

$$\ln L(\lambda) = r n \ln \lambda - n \ln \Gamma(r) + \ln \left(\prod_{i=1}^n y_i^{r-1} \right) - \lambda \sum_{i=1}^n y_i.$$

$$\frac{d}{d\lambda} \ln L(\lambda) = \frac{r n}{\lambda} - \sum_{i=1}^n y_i.$$

$$\frac{d}{d\lambda} \ln L(\lambda) = 0 \quad \Rightarrow \quad \boxed{\lambda_e = \frac{r n}{\sum_{i=1}^n y_i} = \frac{r}{\bar{y}}}.$$

Comment:

- When $r = 1$, this reduces to the exponential distribution case.
- If r is also unknown, it will be much more complicated.
No closed-form solution. One needs numerical solver².
Try MME instead.

²[DW, Example 7.2.25]

A detailed study with data:

E.g. 4. Geometric distribution: $p_X(k; p) = (1 - p)^{k-1} p$, $k = 1, 2, \dots$.

$$L(p) = \prod_{i=1}^n (1 - p)^{k_i - 1} p = (1 - p)^{-n + \sum_{i=1}^n k_i} p^n.$$

$$\ln L(p) = \left(-n + \sum_{i=1}^n k_i \right) \ln(1 - p) + n \ln p.$$

$$\frac{d}{dp} \ln L(p) = -\frac{-n + \sum_{i=1}^n k_i}{1 - p} + \frac{n}{p}.$$

$$\frac{d}{dp} \ln L(p) = 0 \quad \Longrightarrow \quad \boxed{p_e = \frac{n}{\sum_{i=1}^n k_i} = \frac{1}{\bar{k}}}.$$

Comment: Its cousin distribution, the negative binomial distribution can be worked out similarly (See Ex 5.2.14).

k	Observed frequency	Predicted frequency
1	72	74.14
2	35	31.2
3	11	13.13
4	6	5.52
5	2	2.32
6	2	0.98

```

1 # The example from the book.
2 library (pracma) # Load the library "Practical Numerical Math Functions"
3 k<-c(72, 35, 11, 6, 2, 2) # observed freq.
4 a=1:6
5 pe=sum(k)/dot(k,a) # MLE for p.
6 f=a
7 for (i in 1:6) {
8   f[i] = round((1-pe)^(i-1) * pe * sum(k),2)
9 }
10 # Initialize the table
11 d <-matrix(1:18, nrow = 6, ncol = 3)
12 # Now adding the column names
13 colnames(d) <- c("k",
14                  "Observed freq.",
15                  "Predicted freq.")
16 d[1:6,1]<-a
17 d[1:6,2]<-k
18 d[1:6,3]<-f
19 grid.table(d) # Show the table
20 PlotResults("unknown", pe, d, "Geometric.pdf") # Output the results using a user defined function

```


k	Observed frequency	Predicted frequency
1	42	40.96
2	31	27.85
3	15	18.94
4	11	12.88
5	9	8.76
6	5	5.96
7	7	4.05
8	2	2.75
9	1	1.87
10	2	1.27
11	1	0.87
13	1	0.59
14	1	0.4

```

1 # Now let's generate random samples from a Geometric distribution with  $p=1/3$  with the same size
  of the sample.
2 p = 1/3
3 n = 128
4 gdata<-rgeom(n, p)+1 # Generate random samples
5 g<- table(gdata) # Count frequency of your data.
6 g<- t(rbind(as.numeric(rownames(g)), g)) # Transpose and combine two columns.
7 pe=n/dot(g[,1],g[,2]) # MLE for p.
8 f <- g[,1] # Initialize f
9 for (i in 1:nrow(g)) {
10   f[i] = round((1-pe)^(i-1) * pe * n,2)
11 } # Compute the expected frequency
12 g<-cbind(g,f) # Add one columns to your matrix.
13 colnames(g) <- c("k",
14                  "Observed freq.",
15                  "Predicted freq.") # Specify the column names.
16 d_df <- as.data.frame(d) # One can use data frame to store data
17 d_df # Show data on your terminal
18 PlotResults(p, pe, g, "Geometric2.pdf") # Output the results using a user defined function

```

k	Observed frequency	Predicted frequency
1	99	105.88
2	69	68.51
3	47	44.33
4	28	28.69
5	27	18.56
6	9	12.01
7	8	7.77
8	5	5.03
9	5	3.25
10	3	2.11

In case we have several parameters:

E.g. 5. Normal distribution: $f_Y(y; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(y-\mu)^2}{2\sigma^2}}$, $y \in \mathbb{R}$.

$$L(\mu, \sigma^2) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(y_i-\mu)^2}{2\sigma^2}} = (2\pi\sigma^2)^{-n/2} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - \mu)^2\right)$$

$$\ln L(\mu, \sigma^2) = -\frac{n}{2} \ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - \mu)^2.$$

$$\begin{cases} \frac{\partial}{\partial \mu} \ln L(\mu, \sigma^2) = \frac{1}{\sigma^2} \sum_{i=1}^n (y_i - \mu) \\ \frac{\partial}{\partial \sigma^2} \ln L(\mu, \sigma^2) = -\frac{n}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{i=1}^n (y_i - \mu)^2 \end{cases}$$

$$\begin{cases} \frac{\partial}{\partial \mu} \ln L(\mu, \sigma^2) = 0 \\ \frac{\partial}{\partial \sigma^2} \ln L(\mu, \sigma^2) = 0 \end{cases}$$

\Rightarrow

$$\begin{cases} \mu_e = \bar{y} \\ \sigma_e^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})^2 \end{cases}$$

In case when the parameters determine the support of the density:
(Non regular case)

E.g. 6. Uniform distribution on $[a, b]$ with $a < b$: $f_Y(y; a, b) = \frac{1}{b-a}$ if $y \in [a, b]$.

$$L(a, b) = \begin{cases} \prod_{i=1}^n \frac{1}{b-a} = \frac{1}{(b-a)^n} & \text{if } a \leq y_1, \dots, y_n \leq b, \\ 0 & \text{otherwise.} \end{cases}$$

$L(a, b)$ is monotone increasing in a and decreasing in b . Hence, in order to maximize $L(a, b)$, one needs to choose

$$a_e = y_{\min} \quad \text{and} \quad b_e = y_{\max}.$$

E.g. 7. $f_Y(y; \theta) = \frac{2y}{\theta^2}$ for $y \in [0, \theta]$.

$$L(\theta) = \begin{cases} \prod_{i=1}^n \frac{2y_i}{\theta^2} = 2^n \theta^{-2n} \prod_{i=1}^n y_i & \text{if } 0 \leq y_1, \dots, y_n \leq \theta, \\ 0 & \text{otherwise.} \end{cases}$$

\Downarrow

$$\theta_e = y_{\max}.$$

In case of discrete parameter:


E.g. 8. **Wildlife sampling.** Capture-tag-recapture.... In the history, a tags have been put. In order to estimate the population size N , one randomly captures n animals, and there are k tagged. Find the MLE for N .

Sol. The population follows hypergeometric distr.: $p_X(k; N) = \frac{\binom{a}{k} \binom{N-a}{n-k}}{\binom{N}{n}}$.

$$L(N) = \frac{\binom{a}{k} \binom{N-a}{n-k}}{\binom{N}{n}}$$

How to maximize $L(N)$?

```
1 > a=10
2 > k=5
3 > n=20
4 > N=seq(a,a+100)
5 > p=choose(a,k)*choose(N-a,n-k)/
  choose(N,n)
6 > plot(N,p,type = "p")
7 > print(paste("The MLE is", n*a/k))
8 [1] "The MLE is 40"
```



The graph suggests to study the following quantity:

$$r(N) := \frac{L(N)}{L(N-1)} = \frac{N-n}{N} \times \frac{N-a}{N-a-n+k}$$

$$r(N) < 1 \iff na < Nk \text{ i.e., } N > \frac{na}{k}$$

$$N_e = \arg \max \left\{ L(N) : N = \left\lfloor \frac{na}{k} \right\rfloor, \left\lceil \frac{na}{k} \right\rceil \right\}.$$



Method of Moments Estimation

Rationale: The population moments should be close to the sample moments, i.e.,

$$\mathbb{E}(Y^k) \approx \frac{1}{n} \sum_{i=1}^n y_i^k, \quad k = 1, 2, 3, \dots$$

Definition 5.2.3. For a random sample of size n from the discrete (resp. continuous) population/pdf $p_X(k; \theta_1, \dots, \theta_s)$ (resp. $f_Y(y; \theta_1, \dots, \theta_s)$), solutions to

$$\begin{cases} \mathbb{E}(Y) = \frac{1}{n} \sum_{i=1}^n y_i \\ \vdots \\ \mathbb{E}(Y^s) = \frac{1}{n} \sum_{i=1}^n y_i^s \end{cases}$$

which are denoted by $\theta_{1e}, \dots, \theta_{se}$, are called the **method of moments estimates** of $\theta_1, \dots, \theta_s$.

Examples for MME

MME is often the same as MLE:

E.g. 1. Normal distribution: $f_Y(y; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(y-\mu)^2}{2\sigma^2}}$, $y \in \mathbb{R}$.

$$\left\{ \begin{array}{l} \mu = \mathbb{E}(Y) = \frac{1}{n} \sum_{i=1}^n y_i = \bar{y} \\ \sigma^2 + \mu^2 = \mathbb{E}(Y^2) = \frac{1}{n} \sum_{i=1}^n y_i^2 \end{array} \right. \Rightarrow \left\{ \begin{array}{l} \mu_e = \bar{y} \\ \sigma_e^2 = \frac{1}{n} \sum_{i=1}^n y_i^2 - \mu_e^2 \\ \quad = \frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})^2 \end{array} \right.$$

More examples when MLE coincides with MME: Poisson, Exponential, Geometric.

MME is often much more tractable than MLE:

E.g. 2. Gamma distribution³: $f_Y(y; r, \lambda) = \frac{\lambda^r}{\Gamma(r)} y^{r-1} e^{-\lambda y}$ for $y \geq 0$.

$$\begin{cases} \frac{r}{\lambda} = \mathbb{E}(Y) = \frac{1}{n} \sum_{i=1}^n y_i = \bar{y} \\ \frac{r}{\lambda^2} + \frac{r^2}{\lambda^2} = \mathbb{E}(Y^2) = \frac{1}{n} \sum_{i=1}^n y_i^2 \end{cases} \Rightarrow \begin{cases} r_e = \frac{\bar{y}^2}{\hat{\sigma}^2} \\ \lambda_e = \frac{\bar{y}}{\hat{\sigma}^2} = \frac{r_e}{\bar{y}} \end{cases}$$

where \bar{y} is the sample mean and $\hat{\sigma}^2$ is the sample variance:
 $\hat{\sigma}^2 := \frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})^2$.

Comments: MME for λ is consistent with MLE when r is known.

³Check Theorem 4.6.3 on p. 269 for mean and variance

Another tractable example for MME, while less tractable for MLE:


E.g. 3. Neg. binomial distribution: $p_X(k; p, r) = \binom{k+r-1}{k} (1-p)^k p^r$, $k = 0, 1, \dots$.

$$\begin{cases} \frac{r(1-p)}{p} = \mathbb{E}(X) = \bar{k} \\ \frac{r(1-p)}{p^2} = \text{Var}(X) = \hat{\sigma}^2 \end{cases} \Rightarrow \begin{cases} p_e = \frac{\bar{k}}{\hat{\sigma}^2} \\ r_e = \frac{\bar{k}^2}{\hat{\sigma}^2 - \bar{k}} \end{cases}$$

Table 5.2.4

Number	Observed Frequency	Expected Frequency
0-5	0	0
6-10	10	7.7
11-15	20	21.4
16-20	23	28.4
21-25	27	22.4
26-30	12	12.3
31-35	5	5.3
36-40	2	1.8
> 40	1	0.7

Data from: <http://www.seattlecentral.edu/qelp/sets/039/039.html>



$$r_e = 12.74 \text{ and } p_e = 0.391.$$

E.g. 4. $f_Y(y; \theta) = \frac{2y}{\theta^2}$ for $y \in [0, \theta]$.

$$\bar{y} = \mathbb{E}[Y] = \int_0^\theta \frac{2y^2}{\theta^2} dy = \frac{2}{3} \frac{y^3}{\theta^2} \Big|_{y=0}^{y=\theta} = \frac{2}{3} \theta.$$

\Downarrow

$$\boxed{\theta_e = \frac{3}{2} \bar{y}.$$

Chapter 5. Estimation

§ 5.1 Introduction

§ 5.2 Estimating parameters: MLE and MME

§ 5.3 Interval Estimation

§ 5.4 Properties of Estimators

§ 5.5 Minimum-Variance Estimators: The Cramér-Rao Lower Bound

§ 5.6 Sufficient Estimators

§ 5.7 Consistency

§ 5.8 Bayesian Estimation

§ 5.3 Interval Estimation

Rationale. Point estimate doesn't provide precision information.

By using the variance of the estimator, one can construct an interval such that with a high probability that interval will contain the unknown parameter.

- ▶ The interval is called **confidence interval**.
- ▶ The high probability is **confidence level**.

E.g. 1. A random sample of size 4, ($Y_1 = 6.5$, $Y_2 = 9.2$, $Y_3 = 9.9$, $Y_4 = 12.4$), from a normal population:


$$f_Y(y; \mu) = \frac{1}{\sqrt{2\pi} \cdot 0.8} e^{-\frac{1}{2} \left(\frac{y - \mu}{0.8} \right)^2}.$$

Both MLE and MME give $\mu_e = \bar{y} = \frac{1}{4}(6.5 + 9.2 + 9.9 + 12.4) = 9.5$. The estimator $\hat{\mu} = \bar{Y}$ follows normal distribution.

Construct 95%-confidence interval for μ ...

"The parameter is an unknown constant and no probability statement concerning its value may be made."

—Jerzy Neyman, original developer of confidence intervals.



In general, for a normal population with σ known, the $100(1 - \alpha)\%$ **confidence interval** for μ is

$$\left(\bar{y} - z_{\alpha/2} \frac{\sigma}{\sqrt{n}}, \bar{y} + z_{\alpha/2} \frac{\sigma}{\sqrt{n}} \right)$$

Comment: There are many variations

1. One-sided interval such as

$$\left(\bar{y} - z_{\alpha} \frac{\sigma}{\sqrt{n}}, \bar{y} \right) \quad \text{or} \quad \left(\bar{y}, \bar{y} + z_{\alpha} \frac{\sigma}{\sqrt{n}} \right)$$

- | | |
|--|-------------------------------|
| 2. σ is unknown and sample size is small: | z-score \rightarrow t-score |
| 3. σ is unknown and sample size is large: | z-score by CLT |
| 4. Non-Gaussian population but sample size is large: | z-score by CLT |

Theorem. Let k be the number of successes in n independent trials, where n is large and $p = \mathbb{P}(\text{success})$ is unknown. An approximate $100(1 - \alpha)\%$ confidence interval for p is the set of numbers

$$\left(\frac{k}{n} - z_{\alpha/2} \sqrt{\frac{(k/n)(1 - k/n)}{n}}, \frac{k}{n} + z_{\alpha/2} \sqrt{\frac{(k/n)(1 - k/n)}{n}} \right).$$

Proof: It follows the following facts:

- $X \sim \text{binomial}(n, p)$ iff $X = Y_1 + \cdots + Y_n$, while Y_i are i.i.d. Bernoulli(p):

$$\mathbb{E}[Y_i] = p \quad \text{and} \quad \text{Var}(Y_i) = p(1 - p).$$

- **Central Limit Theorem:** Let W_1, W_2, \dots, W_n be an sequence of i.i.d. random variables, whose distribution has mean μ and variance σ^2 , then

$$\frac{\sum_{i=1}^n W_i - n\mu}{\sqrt{n\sigma^2}} \quad \text{approximately follows} \quad N(0, 1), \quad \text{when } n \text{ is large.}$$

- When the sample size n is large, by the central limit theorem,

$$\frac{\sum_{i=1}^n Y_i - np}{\sqrt{np(1-p)}} \stackrel{\text{ap.}}{\approx} N(0, 1)$$

||

$$\frac{X - np}{\sqrt{np(1-p)}} = \frac{\frac{X}{n} - p}{\sqrt{\frac{p(1-p)}{n}}} \approx \frac{\frac{X}{n} - p}{\sqrt{\frac{p_e(1-p_e)}{n}}}$$

- Since $p_e = \frac{k}{n}$, we see that

$$\mathbb{P} \left(-z_{\alpha/2} \leq \frac{\frac{X}{n} - p}{\sqrt{\frac{\frac{k}{n}(1-\frac{k}{n})}{n}}} \leq z_{\alpha/2} \right) \approx 1 - \alpha$$

i.e., the $100(1 - \alpha)\%$ confidence interval for p is

$$\left(\frac{k}{n} - z_{\alpha/2} \sqrt{\frac{(k/n)(1-k/n)}{n}}, \frac{k}{n} + z_{\alpha/2} \sqrt{\frac{(k/n)(1-k/n)}{n}} \right).$$

E.g. 1. Use *median test* to check the randomness of a random generator.

Suppose y_1, \dots, y_n denote measurements presumed to have come from a continuous pdf $f_Y(y)$. Let k denote the number of y_i 's that are less than the median of $f_Y(y)$. If the sample is random, we would expect the difference between $\frac{k}{n}$ and $\frac{1}{2}$ to be small. More specifically, a 95% confidence interval based on k should contain the value 0.5.

Let $f_Y(y) = e^{-y}$. The median is $m = 0.69315$.

```

1  #!/usr/bin/Rscript
2  main <- function() {
3    args <- commandArgs(trailingOnly = TRUE)
4    n <- 100 # Number of random samples.
5    r <- as.numeric(args[1]) # Rate of the exponential
6    # Check if the rate argument is given.
7    if (is.na(r)) return("Please provide the rate and try again.")
8
9    # Now start computing ...
10   f <- function (y) pexp(y, rate = r)-0.5
11   m <- uniroot(f, lower = 0, upper = 100, tol = 1e-9)$root
12   print(paste("For rate ", r, "exponential distribution ",
13             "the median is equal to ", round(m,3)))
14   data <- rexp(n,r) # Generate n random samples
15   data <- round(data,3) # Round to 3 digits after decimal
16   data <- matrix(data, nrow = 10, ncol = 10) # Turn the data to a matrix
17   prmatrix(data) # Show data on terminal
18   k <- sum(data > m) # Count how many entries is bigger than m
19   lowerbd = k/n - 1.96 * sqrt((k/n)*(1-k/n)/n);
20   upperbd = k/n + 1.96 * sqrt((k/n)*(1-k/n)/n);
21   print(paste("The 95% confidence interval is (",
22             round(lowerbd,3), ", ",
23             round(upperbd,3), ")"))
24 }
25 main()

```

Try commandline ...

Math362:./Example-5-3-2.R 1

```
[1] "For rate 1 exponential distribution, the median is equal to 0.693"
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
[1,] 1.324 1.211 0.561 0.640 2.816 2.348 0.788 2.243 1.759 0.103
[2,] 0.476 2.288 0.106 0.079 0.636 1.941 0.801 3.838 0.612 0.030
[3,] 1.085 0.305 0.354 1.013 0.687 1.656 1.043 0.389 1.476 2.158
[4,] 1.267 1.031 0.917 0.681 0.912 0.236 0.054 0.862 0.065 0.402
[5,] 0.957 1.003 1.665 1.137 0.378 1.182 0.659 1.923 1.127 0.364
[6,] 0.307 0.127 0.203 0.394 1.392 2.378 4.192 0.365 3.227 0.337
[7,] 0.707 0.049 0.391 1.967 1.220 2.605 0.887 1.749 1.479 1.526
[8,] 0.662 0.141 0.318 0.523 0.646 1.202 0.442 0.174 1.178 0.177
[9,] 0.397 0.493 0.214 0.522 2.024 4.109 1.268 1.041 0.948 0.382
[10,] 2.260 0.292 0.437 0.962 0.224 4.221 0.594 0.218 0.601 0.941
[1] "The 95% confidence interval is ( 0.422 , 0.618 )"
```

Math362:./Example-5-3-2.R 10

```
[1] "For rate 10 exponential distribution, the median is equal to 0.069"
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
[1,] 0.199 0.069 0.013 0.025 0.000 0.107 0.068 0.116 0.066 0.146
[2,] 0.027 0.076 0.044 0.458 0.052 0.127 0.100 0.100 0.014 0.061
[3,] 0.014 0.078 0.044 0.072 0.028 0.141 0.038 0.022 0.037 0.093
[4,] 0.042 0.015 0.250 0.132 0.292 0.072 0.105 0.244 0.046 0.054
[5,] 0.134 0.074 0.182 0.057 0.021 0.038 0.095 0.196 0.004 0.048
[6,] 0.016 0.021 0.163 0.030 0.139 0.063 0.054 0.006 0.023 0.051
[7,] 0.227 0.055 0.091 0.121 0.066 0.114 0.004 0.021 0.035 0.211
[8,] 0.113 0.083 0.129 0.338 0.160 0.008 0.014 0.167 0.050 0.127
[9,] 0.053 0.073 0.054 0.098 0.004 0.036 0.274 0.276 0.004 0.159
[10,] 0.045 0.469 0.152 0.003 0.129 0.017 0.084 0.072 0.162 0.007
[1] "The 95% confidence interval is ( 0.392 , 0.588 )"
```

Math362:█

Instead of the C.I. $\left(\frac{k}{n} - z_{\alpha/2} \sqrt{\frac{(k/n)(1-k/n)}{n}}, \frac{k}{n} + z_{\alpha/2} \sqrt{\frac{(k/n)(1-k/n)}{n}} \right)$.

One can simply specify the mean $\frac{k}{n}$ and

the **margin of error**: $d := z_{\alpha/2} \sqrt{\frac{(k/n)(1-k/n)}{n}}$.

$$\max_{p \in (0,1)} p(1-p) = p(1-p) \Big|_{p=1/2} = 1/4 \implies d \leq \frac{z_{\alpha/2}}{2\sqrt{n}} =: d_m.$$

Comment:

1. When p is close to $1/2$, $d \approx \frac{z_{\alpha/2}}{2\sqrt{n}}$, which is equivalent to $\sigma_p \approx \frac{1}{2\sqrt{n}}$.

E.g., $n = 1000$, $k/n = 0.48$, and $\alpha = 5\%$, then

$$d = 1.96 \sqrt{\frac{0.48 \times 0.52}{1000}} = 0.0309\bar{7} \quad \text{and} \quad d_m = \frac{1.96}{2\sqrt{1000}} = 0.0309\bar{9}$$

$$\sigma_p = \sqrt{\frac{0.48 \times 0.52}{1000}} = 0.01579873 \quad \text{and} \quad \sigma_p \approx \frac{1}{2\sqrt{1000}} = 0.01581139.$$

2. When p is away from $1/2$, the discrepancy between d and d_m becomes big....

E.g. Running for presidency. Max and Sirius obtained 480 and 520 votes, respectively. What is probability that Max will win?

What if the sample size is $n = 5000$, and Max obtained 2400 votes.

Choosing sample sizes

$$d \leq z_{\alpha/2} \sqrt{p(1-p)/n} \iff n \geq \frac{z_{\alpha/2}^2 p(1-p)}{d^2} \quad (\text{When } p \text{ is known})$$

$$d \leq \frac{z_{\alpha/2}}{2\sqrt{n}} \iff n \geq \frac{z_{\alpha/2}^2}{4d^2} \quad (\text{When } p \text{ is unknown})$$

E.g. Anti-smoking campaign. Need to find an 95% C.I. with a margin of error equal to 1%. Determine the sample size?

$$\text{Answer: } n \geq \frac{1.96^2}{4 \times 0.01^2} = 9640.$$

E.g.' In order to reduce the sample size, a small sample is used to determine p . One finds that $p \approx 0.22$. Determine the sample size again.

$$\text{Answer: } n \geq \frac{1.96^2 \times 0.22 \times 0.78}{\times 0.01^2} = 6592.2.$$

Chapter 5. Estimation

§ 5.1 Introduction

§ 5.2 Estimating parameters: MLE and MME

§ 5.3 Interval Estimation

§ 5.4 Properties of Estimators

§ 5.5 Minimum-Variance Estimators: The Cramér-Rao Lower Bound

§ 5.6 Sufficient Estimators

§ 5.7 Consistency

§ 5.8 Bayesian Estimation

§ 5.4 Properties of Estimators

Question: Estimators are not in general unique (MLE or MME ...). How to select one estimator?



Recall: For a random sample of size n from the population with given pdf, we have X_1, \dots, X_n , which are i.i.d. r.v.'s. The estimator $\hat{\theta}$ is a function of X_i 's:

$$\hat{\theta} = \hat{\theta}(X_1, \dots, X_n).$$

Criteria:

1. Unbiased. (Mean)
2. Efficiency, the minimum-variance estimator. (Variance)
3. Sufficiency.
4. Consistency. (Asymptotic behavior)

Unbiasedness



Definition 5.4.1. Given a random sample of size n whose population distribution depends on an unknown parameter θ , let $\hat{\theta}$ be an estimator of θ .

Then $\hat{\theta}$ is called **unbiased** if $\mathbb{E}(\hat{\theta}) = \theta$;

and $\hat{\theta}$ is called **asymptotically unbiased** if $\lim_{n \rightarrow \infty} \mathbb{E}(\hat{\theta}) = \theta$.

E.g. 1. $f_Y(y; \theta) = \frac{2y}{\theta^2}$ if $y \in [0, \theta]$.

$$- \hat{\theta}_1 = \frac{3}{2} \bar{Y}$$

$$- \hat{\theta}_2 = Y_{\max}.$$

$$- \hat{\theta}_3 = \frac{2n+1}{2n} Y_{\max}.$$

E.g. 2. Let X_1, \dots, X_n be a random sample of size n with the unknown parameter $\theta = \mathbb{E}(X)$. Show that for any constants a_i 's,

$$\hat{\theta} = \sum_{i=1}^n a_i X_i \quad \text{is unbiased} \quad \Longleftrightarrow \quad \sum_{i=1}^n a_i = 1.$$

E.g. 3. Let X_1, \dots, X_n be a random sample of size n with the unknown parameter $\sigma^2 = \text{Var}(X)$.

$$- \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$$

$$- S^2 = \text{Sample Variance} = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$$

$$- S = \text{Sample Standard Deviation} = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2}. \quad (\text{Biased for } \sigma!)$$

E.g. 4. Exponential distr.: $f_Y(y; \lambda) = \lambda e^{-\lambda y}$ for $y \geq 0$. $\hat{\lambda} = 1/\bar{Y}$ is biased.

$n\bar{Y} = \sum_{i=1}^n Y_i \sim \text{Gamma distribution}(n, \lambda)$. Hence,

$$\begin{aligned}\mathbb{E}(\hat{\lambda}) &= \mathbb{E}(1/\bar{Y}) = n \int_0^{\infty} \frac{1}{y} \frac{\lambda^n}{\Gamma(n)} y^{n-1} e^{-\lambda y} dy \\ &= \frac{n\lambda}{n-1} \int_0^{\infty} \underbrace{\frac{\lambda^{n-1}}{\Gamma(n-1)} y^{(n-1)-1} e^{-\lambda y}}_{\text{pdf for Gamma distr. } (n-1, \lambda)} dy \\ &= \frac{n}{n-1} \lambda.\end{aligned}$$


Biased! But $\mathbb{E}(\hat{\lambda}) = \frac{n}{n-1} \lambda \rightarrow \lambda$ as $n \rightarrow \infty$. (Asymptotically unbiased.)

Note: $\hat{\lambda}^* = \frac{n-1}{n\bar{Y}}$ is unbiased.

E.g. 4'. Exponential distr.: $f_Y(y; \theta) = \frac{1}{\theta} e^{-y/\theta}$ for $y \geq 0$. $\hat{\theta} = \bar{Y}$ is unbiased.

$$\mathbb{E}(\hat{\theta}) = \frac{1}{n} \sum_{i=1}^n \mathbb{E}(Y_i) = \frac{1}{n} \sum_{i=1}^n \theta = \theta.$$

Efficiency



Definition 5.4.2. Let $\hat{\theta}_1$ and $\hat{\theta}_2$ be two unbiased estimators for a parameter θ . If $\text{Var}(\hat{\theta}_1) < \text{Var}(\hat{\theta}_2)$, then we say that $\hat{\theta}_1$ is **more efficient** than $\hat{\theta}_2$. The **relative efficiency** of $\hat{\theta}_1$ w.r.t. $\hat{\theta}_2$ is the ratio $\text{Var}(\hat{\theta}_2)/\text{Var}(\hat{\theta}_1)$.

E.g. 1. $f_Y(y; \theta) = \frac{2y}{\theta^2}$ if $y \in [0, \theta]$. Which is more efficient? Find the relative efficiency of $\hat{\theta}_1$ w.r.t. $\hat{\theta}_3$.

$$- \hat{\theta}_1 = \frac{3}{2} \bar{Y}$$

$$- \hat{\theta}_3 = \frac{2n+1}{2n} Y_{\max}.$$

E.g. 2. Let X_1, \dots, X_n be a random sample of size n with the unknown parameter $\theta = \mathbb{E}(X)$ (suppose $\sigma^2 = \text{Var}(X) < \infty$).

Among all possible unbiased estimators $\hat{\theta} = \sum_{i=1}^n a_i X_i$ with $\sum_{i=1}^n a_i = 1$. Find the most efficient one.

Sol:

$$\text{Var}(\hat{\theta}) = \sum_{i=1}^n a_i^2 \text{Var}(X) = \sigma^2 \sum_{i=1}^n a_i^2 \geq \sigma^2 \frac{1}{n} \left(\sum_{i=1}^n a_i \right)^2 = \frac{1}{n} \sigma^2,$$

with equality iff $a_1 = \dots = a_n = 1/n$.

Hence, the most efficient one is the sample mean $\hat{\theta} = \bar{X}$.

□

Chapter 5. Estimation

§ 5.1 Introduction

§ 5.2 Estimating parameters: MLE and MME

§ 5.3 Interval Estimation

§ 5.4 Properties of Estimators

§ 5.5 Minimum-Variance Estimators: The Cramér-Rao Lower Bound

§ 5.6 Sufficient Estimators

§ 5.7 Consistency

§ 5.8 Bayesian Estimation

§ 5.5 MVE: The Cramér-Rao Lower Bound

Question: Can one identify the unbiased estimator having the *smallest* variance?

Short answer: In many cases, yes!

We are going to develop the theory to answer this question in details!

Regular Estimation/Condition: The set of y (resp. k) values, where $f_Y(y; \theta) \neq 0$ (resp. $p_X(k; \theta) \neq 0$), does not depend on θ .

i.e., the domain of the pdf does not depend on the parameter (so that one can differentiate under integration).

Definition. The **Fisher's Information** of a continuous (resp. discrete) random variable Y (resp. X) with pdf $f_Y(y; \theta)$ (resp. $p_X(k; \theta)$) is defined as

$$I(\theta) = \mathbb{E} \left[\left(\frac{\partial \ln f_Y(Y; \theta)}{\partial \theta} \right)^2 \right] \quad \left(\text{resp.} \quad \mathbb{E} \left[\left(\frac{\partial \ln p_X(X; \theta)}{\partial \theta} \right)^2 \right] \right).$$

Lemma. Under regular condition, let Y_1, \dots, Y_n be a random sample of size n from the continuous population pdf $f_Y(y; \theta)$. Then the Fisher Information in the random sample Y_1, \dots, Y_n equals n times the Fisher information in X :

$$\mathbb{E} \left[\left(\frac{\partial \ln f_{Y_1, \dots, Y_n}(Y_1, \dots, Y_n; \theta)}{\partial \theta} \right)^2 \right] = n \mathbb{E} \left[\left(\frac{\partial \ln f_Y(Y; \theta)}{\partial \theta} \right)^2 \right] = n I(\theta). \quad (1)$$

(A similar statement holds for the discrete case $p_X(k; \theta)$).

Proof. Based on two observations:

$$\begin{aligned} LHS &= \mathbb{E} \left[\left(\sum_{i=1}^n \frac{\partial}{\partial \theta} \ln f_{Y_i}(Y_i; \theta) \right)^2 \right] \\ \mathbb{E} \left(\frac{\partial}{\partial \theta} \ln f_{Y_i}(Y_i; \theta) \right) &= \int_{\mathbb{R}} \frac{\frac{\partial}{\partial \theta} f_Y(y; \theta)}{f_Y(y; \theta)} f_Y(y; \theta) dy = \int_{\mathbb{R}} \frac{\partial}{\partial \theta} f_Y(y; \theta) dy \\ &\stackrel{\text{R.C.}}{=} \frac{\partial}{\partial \theta} \int_{\mathbb{R}} f_Y(y; \theta) dy = \frac{\partial}{\partial \theta} 1 = 0. \end{aligned}$$

□

Lemma. Under regular condition, if $\ln f_Y(y; \theta)$ is twice differentiable in θ , then

$$I(\theta) = -\mathbb{E} \left[\frac{\partial^2}{\partial \theta^2} \ln f_Y(Y; \theta) \right]. \quad (2)$$

(A similar statement holds for the discrete case $p_X(k; \theta)$).

Proof. This is due to the two facts:

$$\begin{aligned} \frac{\partial^2}{\partial \theta^2} \ln f_Y(Y; \theta) &= \frac{\frac{\partial^2}{\partial \theta^2} f_Y(Y; \theta)}{f_Y(Y; \theta)} - \underbrace{\left(\frac{\frac{\partial}{\partial \theta} f_Y(Y; \theta)}{f_Y(Y; \theta)} \right)^2}_{= \left(\frac{\partial}{\partial \theta} \ln f_Y(Y; \theta) \right)^2} \\ &= \left(\frac{\partial}{\partial \theta} \ln f_Y(Y; \theta) \right)^2 \end{aligned}$$

$$\begin{aligned} \mathbb{E} \left(\frac{\frac{\partial^2}{\partial \theta^2} f_Y(Y; \theta)}{f_Y(Y; \theta)} \right) &= \int_{\mathbb{R}} \frac{\frac{\partial^2}{\partial \theta^2} f_Y(y; \theta)}{f_Y(y; \theta)} f_Y(y; \theta) dy = \int_{\mathbb{R}} \frac{\partial^2}{\partial \theta^2} f_Y(y; \theta) dy. \\ &\stackrel{R.C.}{=} \frac{\partial^2}{\partial \theta^2} \int_{\mathbb{R}} f_Y(y; \theta) dy = \frac{\partial^2}{\partial \theta^2} 1 = 0. \end{aligned}$$

□

Theorem (Cramér-Rao Inequality) Under regular condition, let Y_1, \dots, Y_n be a random sample of size n from the continuous population pdf $f_Y(y; \theta)$. Let $\hat{\theta} = \hat{\theta}(Y_1, \dots, Y_n)$ be any unbiased estimator for θ . Then

$$\text{Var}(\hat{\theta}) \geq \frac{1}{n I(\theta)}.$$

(A similar statement holds for the discrete case $p_X(k; \theta)$).

Proof. If $n = 1$, then by Cauchy-Schwartz inequality,

$$\mathbb{E} \left[(\hat{\theta} - \theta) \frac{\partial}{\partial \theta} \ln f_Y(Y; \theta) \right] \leq \sqrt{\text{Var}(\hat{\theta}) \times I(\theta)}$$

On the other hand,

$$\begin{aligned} \mathbb{E} \left[(\hat{\theta} - \theta) \frac{\partial}{\partial \theta} \ln f_Y(Y; \theta) \right] &= \int_{\mathbb{R}} (\hat{\theta} - \theta) \frac{\frac{\partial}{\partial \theta} f_Y(y; \theta)}{f_Y(y; \theta)} f_Y(y; \theta) dy \\ &= \int_{\mathbb{R}} (\hat{\theta} - \theta) \frac{\partial}{\partial \theta} f_Y(y; \theta) dy \\ &= \frac{\partial}{\partial \theta} \underbrace{\int_{\mathbb{R}} (\hat{\theta} - \theta) f_Y(y; \theta) dy}_{= \mathbb{E}(\hat{\theta} - \theta) = 0} + 1 = 1. \end{aligned}$$

For general n , apply for (1).


□.

Definition. Let Θ be the set of all estimators $\hat{\theta}$ that are unbiased for the parameter θ . We say that $\hat{\theta}^*$ is a **best** or **minimum-variance** estimator (MVE) if $\hat{\theta}^* \in \Theta$ and

$$\text{Var}(\hat{\theta}^*) \leq \text{Var}(\hat{\theta}) \quad \text{for all } \hat{\theta} \in \Theta.$$

Definition. An unbiased estimator $\hat{\theta}$ is **efficient** if $\text{Var}(\hat{\theta})$ is equal to the Cramér-Rao lower bound, i.e., $\text{Var}\hat{\theta} = (n I(\theta))^{-1}$.

The **efficiency** of an unbiased estimator $\hat{\theta}$ is defined to be $\left(n I(\theta) \text{Var}(\hat{\theta}) \right)^{-1}$.



E.g. 1. $X \sim \text{Bernoulli}(p)$. Check whether $\hat{p} = \bar{X}$ is efficient?

Step 1. Compute Fisher's Information:

$$p_X(k; p) = p^k (1 - p)^{1-k}.$$

$$\ln p_X(k; p) = k \ln p + (1 - k) \ln(1 - p)$$

$$\frac{\partial}{\partial p} \ln p_X(k; p) = \frac{k}{p} - \frac{1 - k}{1 - p}$$

$$-\frac{\partial^2}{\partial^2 p} \ln p_X(k; p) = \frac{k}{p^2} + \frac{1 - k}{(1 - p)^2}$$

$$-\mathbb{E} \left[\frac{\partial^2}{\partial^2 p} \ln p_X(X; p) \right] = \mathbb{E} \left[\frac{X}{p^2} + \frac{1 - X}{(1 - p)^2} \right] = \frac{1}{p} + \frac{1}{1 - p} = \frac{1}{pq}.$$

$$I(p) = \frac{1}{pq}, \quad q = 1 - p.$$

Step 2. Compute $\text{Var}(\hat{p})$.

$$\text{Var}(\hat{p}) = \frac{1}{n^2} \text{Var} \left(\sum_{i=1}^n X_i \right) = \frac{1}{n^2} npq = \frac{pq}{n}$$

Conclusion Because \hat{p} is unbiased and $\text{Var}(\hat{p}) = (nI(p))^{-1}$, \hat{p} is efficient.

E.g. 2. Exponential distr.: $f_Y(y; \lambda) = \lambda e^{-\lambda y}$ for $y \geq 0$. Is $\hat{\lambda} = 1/\bar{Y}$ efficient?

Answer No, because $\hat{\lambda}$ is biased. Nevertheless, we can still compute Fisher's Information as follows

Fisher's Inf.

$$\ln f_Y(y; \lambda) = \ln \lambda - \lambda y$$

$$\frac{\partial}{\partial \lambda} \ln f_Y(y; \lambda) = \frac{1}{\lambda} - y$$

$$-\frac{\partial^2}{\partial^2 \lambda} \ln f_Y(y; \lambda) = \frac{1}{\lambda^2}$$

$$-\mathbb{E} \left[\frac{\partial^2}{\partial^2 \lambda} \ln f_Y(Y; \lambda) \right] = \mathbb{E} \left[\frac{1}{\lambda^2} \right] = \frac{1}{\lambda^2}.$$

$$I(\lambda) = \lambda^{-2}$$

Try: $\hat{\lambda}^* := \frac{n-1}{n} \frac{1}{\bar{Y}}$. It is unbiased. Is it efficient?

E.g. 2'. Exponential distr.: $f_Y(y; \theta) = \theta^{-1} e^{-y/\theta}$ for $y \geq 0$. $\hat{\theta} = \bar{Y}$ efficient?

Step. 1. Compute Fisher's Information:

$$\ln f_Y(y; \theta) = -\ln \theta - \frac{y}{\theta}$$

$$\frac{\partial}{\partial \theta} \ln f_Y(y; \theta) = -\frac{1}{\theta} + \frac{y}{\theta^2}$$

$$-\frac{\partial^2}{\partial^2 \theta} \ln f_Y(y; \theta) = -\frac{1}{\theta^2} + \frac{2y}{\theta^3}$$

$$-\mathbb{E} \left[\frac{\partial^2}{\partial^2 \theta} \ln f_Y(Y; \theta) \right] = \mathbb{E} \left[-\frac{1}{\theta^2} + \frac{2Y}{\theta^3} \right] = -\frac{1}{\theta^2} + \frac{2\theta}{\theta^3} = \theta^{-2}.$$

$$\boxed{I(\theta) = \theta^{-2}}$$

Step 2. Compute $\text{Var}(\hat{\theta})$:

$$\text{Var}(\bar{Y}) = \frac{1}{n^2} \sum_{i=1}^n \text{Var}(Y_i) = \frac{1}{n^2} n\theta^2 = \frac{\theta^2}{n}.$$

Conclusion. Because $\hat{\theta}$ is unbiased and $\text{Var}(\hat{\theta}) = (nI(\theta))^{-1}$, $\hat{\theta}$ is efficient.

E.g. 3. $f_Y(y; \theta) = 2y/\theta^2$ for $y \in [0, \theta]$. $\hat{\theta} = \frac{3}{2}\bar{Y}$ efficient?

Step. 1. Compute Fisher's Information:

$$\ln f_Y(y; \theta) = \ln(2y) - 2 \ln \theta$$

$$\frac{\partial}{\partial \theta} \ln f_Y(y; \theta) = -\frac{2}{\theta}$$

By the definition of Fisher's information,

$$I(\theta) = \mathbb{E} \left[\left(\frac{\partial}{\partial \theta} \ln f_Y(y; \theta) \right)^2 \right] = \mathbb{E} \left[\left(-\frac{2}{\theta} \right)^2 \right] = \frac{4}{\theta^2}.$$

However, if we compute

$$-\frac{\partial^2}{\partial^2 \theta} \ln f_Y(y; \theta) = -\frac{2}{\theta^2}$$

$$-\mathbb{E} \left[\frac{\partial^2}{\partial^2 \theta} \ln f_Y(Y; \theta) \right] = \mathbb{E} \left[-\frac{2}{\theta^2} \right] = -\frac{2}{\theta^2} \neq \frac{4}{\theta^2} = I(\theta). \quad (\dagger)$$

Step 2. Compute $\text{Var}(\hat{\theta})$:

$$\text{Var}(\hat{\theta}) = \frac{9}{4n} \text{Var}(Y) = \frac{9}{4n} \frac{\theta^2}{18} = \frac{\theta^2}{8n}.$$

Discussion. Even though $\hat{\theta}$ is unbiased, we have two discrepancies: (\dagger) and

$$\text{Var}(\hat{\theta}) = \frac{\theta^2}{8n} \leq \frac{\theta^2}{4n} = \frac{1}{nI(\theta)}$$

This is because this is not a regular estimation!

Chapter 5. Estimation

§ 5.1 Introduction

§ 5.2 Estimating parameters: MLE and MME

§ 5.3 Interval Estimation

§ 5.4 Properties of Estimators

§ 5.5 Minimum-Variance Estimators: The Cramér-Rao Lower Bound

§ 5.6 Sufficient Estimators

§ 5.7 Consistency

§ 5.8 Bayesian Estimation

§ 5.6 Sufficient Estimators

Rationale: Let $\hat{\theta}$ be an estimator to the unknown parameter θ . Whether does $\hat{\theta}$ contain all information about θ ?

Equivalently, how can one reduce the random sample of size n , denoted by (X_1, \dots, X_n) , to a function without losing any information about θ ?

E.g., let's choose the function $h(X_1, \dots, X_n) := \frac{1}{n} \sum_{i=1}^n X_i$, the sample mean. In many cases, $h(X_1, \dots, X_n)$ contains all relevant information about the true mean $\mathbb{E}(X)$. In that case, $h(X_1, \dots, X_n)$, as an estimator, is sufficient.

Definition. Let (X_1, \dots, X_n) be a random sample of size n from a discrete population with a unknown parameter θ , of which $\hat{\theta}$ (resp. θ_e) be an estimator (resp. estimate). We call $\hat{\theta}$ and θ_e **sufficient** if

$$\mathbb{P} \left(X_1 = k_1, \dots, X_n = k_n \mid \hat{\theta} = \theta_e \right) = b(k_1, \dots, k_n) \quad (\text{Sufficiency-1})$$

is a function that does not depend on θ .

In case for random sample (Y_1, \dots, Y_n) from the continuous population, (Sufficiency-1) should be

$$f_{Y_1, \dots, Y_n \mid \hat{\theta} = \theta_e} \left(y_1, \dots, y_n \mid \hat{\theta} = \theta_e \right) = b(y_1, \dots, y_n)$$

Note: $\hat{\theta} = h(X_1, \dots, X_n)$ and $\theta_e = h(k_1, \dots, k_n)$.
or $\hat{\theta} = h(Y_1, \dots, Y_n)$ and $\theta_e = h(y_1, \dots, y_n)$.

Equivalently,

Definition. ... $\hat{\theta}$ (or θ_e) is **sufficient** if the likelihood function can be factorized as:

$$L(\theta) = \begin{cases} \prod_{i=1}^n p_X(k_i; \theta) = g(\theta_e, \theta) b(k_1, \dots, k_n) & \text{Discrete} \\ \prod_{i=1}^n f_Y(y_i; \theta) = g(\theta_e, \theta) b(y_1, \dots, y_n) & \text{Continuous} \end{cases} \quad (\text{Sufficiency-2})$$

where g is a function of two arguments only and b is a function that does not depend on θ .

E.g. 1. A random sample of size n from Bernoulli(P). $\hat{p} = \sum_{i=1}^n X_i$. Check sufficiency of \hat{p} for p by (Sufficiency-1):

Case I: If $k_1, \dots, k_n \in \{0, 1\}$ such that $\sum_{i=1}^n k_i \neq c$, then

$$\mathbb{P}(X_1 = k_1, \dots, X_n = k_n \mid \hat{p} = c) = 0.$$

Case II: If $k_1, \dots, k_n \in \{0, 1\}$ such that $\sum_{i=1}^n k_i = c$, then

$$\begin{aligned}
& \mathbb{P}(X_1 = k_1, \dots, X_n = k_n \mid \hat{p} = c) \\
&= \frac{\mathbb{P}(X_1 = k_1, \dots, X_n = k_n, \hat{p} = c)}{\mathbb{P}(\hat{p} = c)} \\
&= \frac{\mathbb{P}(X_1 = k_1, \dots, X_n = k_n, X_n + \sum_{i=1}^{n-1} X_i = c)}{\mathbb{P}(\sum_{i=1}^n X_i = c)} \\
&= \frac{\mathbb{P}(X_1 = k_1, \dots, X_{n-1} = k_{n-1}, X_n = c - \sum_{i=1}^{n-1} k_i)}{\mathbb{P}(\sum_{i=1}^n X_i = c)} \\
&= \frac{(\prod_{i=1}^{n-1} p^{k_i} (1-p)^{1-k_i}) \times p^{c - \sum_{i=1}^{n-1} k_i} (1-p)^{1-c + \sum_{i=1}^{n-1} k_i}}{\binom{n}{c} p^c (1-p)^{n-c}} \\
&= \frac{1}{\binom{n}{c}}.
\end{aligned}$$

In summary,

$$\mathbb{P}(X_1 = k_1, \dots, X_n = k_n \mid \hat{p} = c) = \begin{cases} \frac{1}{\binom{n}{c}} & \text{if } k_i \in \{0, 1\} \text{ s.t. } \sum_{i=1}^n k_i = c, \\ 0 & \text{otherwise.} \end{cases}$$

Hence, by (Sufficiency-1), $\hat{p} = \sum_{i=1}^n X_i$ is a sufficient estimator for p .

E.g. 1'. As in E.g. 1, check sufficiency of \hat{p} for p by (Sufficiency-2):

Notice that $p_e = \sum_{i=1}^n k_i$. Then

$$\begin{aligned} L(p) &= \prod_{i=1}^n p_X(k_i; p) = \prod_{i=1}^n p^{k_i} (1-p)^{1-k_i} \\ &= p^{\sum_{i=1}^n k_i} (1-p)^{n - \sum_{i=1}^n k_i} \\ &= p^{p_e} (1-p)^{n-p_e} \end{aligned}$$

Therefore, p_e (or \hat{p}) is sufficient since (Sufficiency-2) is satisfied with

$$g(p_e, p) = p^{p_e} (1-p)^{n-p_e} \quad \text{and} \quad b(k_1, \dots, k_n) = 1.$$

Comment 1. The estimator \hat{p} is sufficient but not unbiased since $\mathbb{E}(\hat{p}) = np \neq p$.

2. Any one-to-one function of a sufficient estimator is again a sufficient estimator. E.g., $\hat{p}_2 := \frac{1}{n}\hat{p}$, which is a unbiased, sufficient, and MVE.

3. $\hat{p}_3 := X_1$ is not sufficient!

E.g. 2. Poisson(λ), $p_X(k; \lambda) = e^{-\lambda} \lambda^k / k!$, $k = 0, 1, \dots$. Show that $\hat{\lambda} = (\sum_{i=1}^n X_i)^2$ is sufficient for λ for a sample of size n .

Sol: The Corresponding estimate is $\lambda_e = (\sum_{i=1}^n k_i)^2$.

$$\begin{aligned} L(\lambda) &= \prod_{i=1}^n \frac{e^{-\lambda} \lambda^{k_i}}{k_i!} \\ &= e^{-n\lambda} \lambda^{\sum_{i=1}^n k_i} \left(\prod_{i=1}^n k_i! \right)^{-1} \\ &= \underbrace{e^{-n\lambda} \lambda^{\sqrt{\lambda_e}}}_{g(\lambda_e, \lambda)} \times \underbrace{\left(\prod_{i=1}^n k_i! \right)^{-1}}_{b(k_1, \dots, k_n)}. \end{aligned}$$

Hence, $\hat{\lambda}$ is sufficient estimator for λ .

□

E.g. 3. Let Y_1, \dots, Y_n be a random sample from $f_Y(y; \theta) = \frac{2y}{\theta^2}$ for $y \in [0, \theta]$.
Whether is the MLE $\hat{\theta} = Y_{\max}$ sufficient for θ ?

Sol: The corresponding estimate is $\theta_e = y_{\max}$.

$$\begin{aligned} L(\theta) &= \prod_{i=1}^n \frac{2y}{\theta^2} I_{[0, \theta]}(y_i) = 2^n \theta^{-2n} \left(\prod_{i=1}^n y_i \right) \times \prod_{i=1}^n I_{[0, \theta]}(y_i) \\ &= 2^n \theta^{-2n} \left(\prod_{i=1}^n y_i \right) \times I_{[0, \theta]}(y_{\max}) \\ &= \underbrace{2^n \theta^{-2n} I_{[0, \theta]}(\theta_e)}_{=g(\theta_e, \theta)} \times \underbrace{\prod_{i=1}^n y_i}_{=b(y_1, \dots, y_k)}. \end{aligned}$$

Hence, $\hat{\theta}$ is a sufficient estimator for θ . □

Note: MME $\hat{\theta} = \frac{3}{2} \bar{Y}$ is NOT sufficient for θ !

Chapter 5. Estimation

§ 5.1 Introduction

§ 5.2 Estimating parameters: MLE and MME

§ 5.3 Interval Estimation

§ 5.4 Properties of Estimators

§ 5.5 Minimum-Variance Estimators: The Cramér-Rao Lower Bound

§ 5.6 Sufficient Estimators

§ 5.7 Consistency

§ 5.8 Bayesian Estimation

Definition. An estimator $\hat{\theta}_n = h(W_1, \dots, W_n)$ is said to be **consistent** if it converges to θ *in probability*, i.e., for any $\epsilon > 0$,

$$\lim_{n \rightarrow \infty} \mathbb{P} \left(|\hat{\theta}_n - \theta| < \epsilon \right) = 1.$$

Comment: In the ϵ - δ language, the above **convergence in probability** says

$$\forall \epsilon > 0, \forall \delta > 0, \exists n(\epsilon, \delta) > 0, \text{ s.t. } \forall n \geq n(\epsilon, \delta),$$

$$\mathbb{P} \left(|\hat{\theta}_n - \theta| < \epsilon \right) > 1 - \delta.$$

A useful tool to check convergence in probability is

Theorem. (Chebyshev's inequality) Let W be any r.v. with finite mean μ and variance σ^2 . Then for any $\epsilon > 0$

$$\mathbb{P}(|W - \mu| < \epsilon) \geq 1 - \frac{\sigma^2}{\epsilon^2},$$

or, equivalently,

$$\mathbb{P}(|W - \mu| \geq \epsilon) \leq \frac{\sigma^2}{\epsilon^2}.$$

Proof. ...



As a consequence of Chebyshev's inequality, we have

Proposition. The sample mean $\hat{\mu}_n = \frac{1}{n} \sum_{i=1}^n W_i$ is consistent for $\mathbb{E}(W) = \mu$, provided that the population W has finite mean μ and variance σ^2 .

Proof.

$$\mathbb{E}(\hat{\mu}_n) = \mu \quad \text{and} \quad \text{Var}(\hat{\mu}_n) = \frac{\sigma^2}{n}.$$

$$\forall \epsilon > 0, \quad \mathbb{P}(|\hat{\mu}_n - \mu| \leq \epsilon) \geq 1 - \frac{\sigma^2}{n\epsilon^2} \rightarrow 1.$$

□

E.g. 1. Let Y_1, \dots, Y_n be a random sample of size n from the uniform pdf $f_Y(y; \theta) = 1/\theta$, $y \in [0, \theta]$. Let $\hat{\theta}_n = Y_{\max}$. We know that Y_{\max} is biased. Is it consistent?

Sol. The c.d.f. of Y is equal to $F_Y(y) = y/\theta$ for $y \in [0, \theta]$. Hence,

$$f_{Y_{\max}}(y) = nF_Y(y)^{n-1}f_Y(y) = \frac{ny^{n-1}}{\theta^n}, \quad y \in [0, \theta].$$

Therefore,

$$\begin{aligned}\mathbb{P}(|\hat{\theta}_n - \theta| < \epsilon) &= \mathbb{P}(\theta - \epsilon < \hat{\theta}_n < \theta + \epsilon) \\ &= \int_{\theta-\epsilon}^{\theta} \frac{ny^{n-1}}{\theta^n} dy + \int_{\theta}^{\theta+\epsilon} 0 dy \\ &= 1 - \left(\frac{\theta - \epsilon}{\theta}\right)^n \\ &\rightarrow 0 \quad \text{as } n \rightarrow \infty.\end{aligned}$$

□

E.g. 2. Suppose Y_1, Y_2, \dots, Y_n is a random sample from the exponential pdf, $f_Y(y; \lambda) = \lambda e^{-\lambda y}$, $y > 0$. Show that $\hat{\lambda}_n = Y_1$ is not consistent for λ .

Sol. To prove $\hat{\lambda}_n$ is not consistent for λ , we need only to find out one $\epsilon > 0$ such that the following limit does not hold:

$$\lim_{n \rightarrow \infty} \mathbb{P} \left(|\hat{\lambda}_n - \lambda| < \epsilon \right) = 1. \quad (3)$$

We can choose $\epsilon = \lambda/m$ for any $m \geq 1$. Then

$$\begin{aligned} |\hat{\lambda}_n - \lambda| \leq \frac{\lambda}{m} &\iff \left(1 - \frac{1}{m}\right) \lambda \leq \hat{\lambda}_n \leq \left(1 + \frac{1}{m}\right) \lambda \\ &\implies \hat{\lambda}_n \geq \left(1 - \frac{1}{m}\right) \lambda. \end{aligned}$$

Hence,

$$\begin{aligned}\mathbb{P}\left(|\hat{\lambda}_n - \lambda| < \frac{\lambda}{m}\right) &\leq \mathbb{P}\left(\hat{\lambda}_n \geq \left(1 - \frac{1}{m}\right)\lambda\right) \\ &= \mathbb{P}\left(Y_1 \geq \left(1 - \frac{1}{m}\right)\lambda\right) \\ &= \int_{\left(1 - \frac{1}{m}\right)\lambda}^{\infty} \lambda e^{-\lambda y} dy \\ &= e^{-\left(1 - \frac{1}{m}\right)\lambda^2} < 1.\end{aligned}$$

Therefore, the limit in (3) cannot hold. □

Chapter 5. Estimation

§ 5.1 Introduction

§ 5.2 Estimating parameters: MLE and MME

§ 5.3 Interval Estimation

§ 5.4 Properties of Estimators

§ 5.5 Minimum-Variance Estimators: The Cramér-Rao Lower Bound

§ 5.6 Sufficient Estimators

§ 5.7 Consistency

§ 5.8 Bayesian Estimation

Rationale: Let W be an estimator dependent on a parameter θ .

1. Frequentists view θ as a parameter whose exact value is to be estimated.
2. Bayesians view θ is the value of a random variable Θ .

One can incorporate our knowledge on Θ — the **prior distribution** $p_{\Theta}(\theta)$ if Θ is discrete and $f_{\Theta}(\theta)$ if Θ is continuous — and use Bayes' formula to update our knowledge on Θ upon new observation $W = w$:

$$g_{\Theta}(\theta|W = w) = \begin{cases} \frac{p_W(w|\Theta = \theta)p_{\Theta}(\theta)}{\mathbb{P}(W = w)} & \text{if } W \text{ is discrete} \\ \frac{f_W(w|\Theta = \theta)f_{\Theta}(\theta)}{f_W(w)} & \text{if } W \text{ is continuous} \end{cases}$$

where $g_{\Theta}(\theta|W = w)$ is called **posterior distribution** of Θ .

Likelihood
of sample W

Prior distri-
bution of Θ

$$P(\Theta|W) = \frac{P(W|\Theta)P(\Theta)}{P(W)}$$

Poste-
rior of Θ

Total
Probability
of sample W

Four cases for computing posterior distribution

$g_{\Theta}(\theta W = w)$	W discrete	W continuous
Θ discrete	$\frac{p_W(w \Theta = \theta)p_{\Theta}(\theta)}{\sum_i p_W(w \Theta = \theta_i)p_{\Theta}(\theta_i)}$	$\frac{f_W(w \Theta = \theta)p_{\Theta}(\theta)}{\sum_i f_W(w \Theta = \theta_i)p_{\Theta}(\theta_i)}$
Θ continuous	$\frac{p_W(w \Theta = \theta)f_{\Theta}(\theta)}{\int_{\mathbb{R}} p_W(w \Theta = \theta')f_{\Theta}(\theta')d\theta'}$	$\frac{f_W(w \Theta = \theta)f_{\Theta}(\theta)}{\int_{\mathbb{R}} f_W(w \Theta = \theta')f_{\Theta}(\theta')d\theta'}$

Gamma distributions

$$\Gamma(r) := \int_0^\infty y^{r-1} e^{-y} dy, \quad r > 0.$$

Two parametrizations for **Gamma distributions**:


1. With a **shape parameter** r and a **scale parameter** θ :

$$f_Y(y; r, \theta) = \frac{y^{r-1} e^{-y/\theta}}{\theta^r \Gamma(r)}, \quad y > 0, r, \theta > 0.$$

2. With a **shape parameter** r and a **rate parameter** $\lambda = 1/\theta$,

$$f_Y(y; r, \lambda) = \frac{\lambda^r y^{r-1} e^{-\lambda y}}{\Gamma(r)}, \quad y > 0, r, \lambda > 0.$$

$$\mathbb{E}[Y] = \frac{r}{\lambda} = r\theta \quad \text{and} \quad \text{Var}(Y) = \frac{r}{\lambda^2} = r\theta^2$$



```
1 # Plot gamma distributions
2 x = seq(0,20,0.01)
3 k= 3 # Shape parameter
4 theta = 0.5 # Scale parameter
5 plot(x,dgamma(x, k, scale = theta),
6      type="l",
7      col="red")
```

Beta distributions


$$B(\alpha, \beta) := \int_0^1 y^{\alpha-1} (1-y)^{\beta-1} dy, \quad \alpha, \beta > 0.$$

$$\begin{aligned} & \vdots \quad \vdots \\ &= \frac{\Gamma(\alpha)\Gamma(\beta)}{\Gamma(\alpha+\beta)}. \end{aligned} \quad (\text{see Appendix})$$

Beta distribution

$$f_Y(y; \alpha, \beta) = \frac{y^{\alpha-1} (1-y)^{\beta-1}}{B(\alpha, \beta)}, \quad y \in [0, 1], \alpha, \beta > 0.$$

$$\mathbb{E}[Y] = \frac{\alpha}{\alpha + \beta} \quad \text{and} \quad \text{Var}(Y) = \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)}$$



```

1 # Plot Beta distributions
2 x = seq(0,1,0.01)
3 a = 13
4 b = 2
5 plot(x, dbeta(x,a,b),
6      type="l",
7      col="red")

```

E.g. 1. Let X_1, \dots, X_n be a random sample from Bernoulli(θ):
 $p_{X_i}(k; \theta) = \theta^k(1 - \theta)^{1-k}$ for $k = 0, 1$.

Let $X = \sum_{i=1}^n X_i$. Then X follows binomial(n, θ).

Prior distribution: $\Theta \sim \text{beta}(r, s)$, i.e., $f_{\Theta}(\theta) = \frac{\Gamma(r+s)}{\Gamma(r)\Gamma(s)} \theta^{r-1}(1 - \theta)^{s-1}$ for $\theta \in [0, 1]$.

$$\begin{array}{ll}
 X_1, \dots, X_n \mid \theta & \sim \text{Bernoulli}(\theta) \\
 \Theta & \sim \text{Beta}(r, s) \\
 & r \text{ \& } s \text{ are known}
 \end{array}
 \qquad
 \begin{array}{ll}
 X = \sum_{i=1}^n X_i \mid \theta & \sim \text{Binomial}(n, \theta) \\
 \Theta & \sim \text{Beta}(r, s) \\
 & r \text{ \& } s \text{ are known}
 \end{array}$$

Example
5.8.2

Max, a video game pirate (and Bayesian), is trying to decide how many illegal copies of *Zombie Beach Party* to have on hand for the upcoming holiday season. To get a rough idea of what the demand might be, he talks with n potential customers and finds that $X = k$ would buy a copy for a present (or for themselves). The obvious choice for a probability model for X , of course, would be the binomial pdf. Given n potential customers, the probability that k would actually buy one of Max's illegal copies is the familiar

$$p_X(k | \theta) = \binom{n}{k} \theta^k (1 - \theta)^{n-k}, \quad k = 0, 1, \dots, n$$

where the maximum likelihood estimate for θ is given by $\theta_e = \frac{k}{n}$.

It may very well be the case, though, that Max has some additional insight about the value of θ on the basis of similar video games that he illegally marketed in previous years. Suppose he suspects, for example, that the percentage of potential customers who will buy *Zombie Beach Party* is likely to be between 3% and 4% and probably will not exceed 7%. A reasonable prior distribution for Θ , then, would be a pdf mostly concentrated over the interval 0 to 0.07 with a mean or median in the 0.035 range.

One such probability model whose shape would comply with the restraints that Max is imposing is the *beta pdf*. Written with Θ as the random variable, the (two-parameter) beta pdf is given by

$$f_{\Theta}(\theta) = \frac{\Gamma(r+s)}{\Gamma(r)\Gamma(s)} \theta^{r-1} (1-\theta)^{s-1}, \quad 0 \leq \theta \leq 1$$

The beta distribution with $r = 2$ and $s = 4$ is pictured in Figure 5.8.1. By choosing different values for r and s , $f_{\Theta}(\theta)$ can be skewed more sharply to the right or to the left, and the bulk of the distribution can be concentrated close to zero or close to one. The question is, if an appropriate beta pdf is used as a *prior* distribution for Θ , and if a random sample of k potential customers (out of n) said they would buy the video game, what would be a reasonable *posterior* distribution for Θ ?




Figure 5.8.1

X is discrete and Θ is continuous.

$$g_{\Theta}(\theta|X = k) = \frac{p_X(k|\Theta = \theta)f_{\Theta}(\theta)}{\int_{\mathbb{R}} p_X(k|\Theta = \theta')f_{\Theta}(\theta')d\theta'}$$

$$\begin{aligned} p_X(k|\Theta = \theta)f_{\Theta}(\theta) &= \binom{n}{k} \theta^k (1 - \theta)^{n-k} \times \frac{\Gamma(r + s)}{\Gamma(r)\Gamma(s)} \theta^{r-1} (1 - \theta)^{s-1} \\ &= \binom{n}{k} \frac{\Gamma(r + s)}{\Gamma(r)\Gamma(s)} \theta^{k+r-1} (1 - \theta)^{n-k+s-1}, \quad \theta \in [0, 1]. \end{aligned}$$

$$\begin{aligned} p_X(k) &= \int_{\mathbb{R}} p_X(k|\Theta = \theta')f_{\Theta}(\theta')d\theta' \\ &= \binom{n}{k} \frac{\Gamma(r + s)}{\Gamma(r)\Gamma(s)} \int_0^1 \theta'^{k+r-1} (1 - \theta')^{n-k+s-1} d\theta' \\ &= \binom{n}{k} \frac{\Gamma(r + s)}{\Gamma(r)\Gamma(s)} \times \frac{\Gamma(k + r)\Gamma(n - k + s)}{\Gamma((k + r) + (n - k + s))} \end{aligned}$$

$$\begin{aligned}
 g_{\Theta}(\theta|X = k) &= \frac{\binom{n}{k} \frac{\Gamma(r+s)}{\Gamma(r)\Gamma(s)} \times \theta^{k+r-1} (1-\theta)^{n-k+s-1}}{\binom{n}{k} \frac{\Gamma(r+s)}{\Gamma(r)\Gamma(s)} \times \frac{\Gamma(k+r)\Gamma(n-k+s)}{\Gamma((k+r) + (n-k+s))}} \\
 &= \frac{\Gamma(n+r+s)}{\Gamma(k+r)\Gamma(n-k+s)} \theta^{k+r-1} (1-\theta)^{n-k+s-1}, \quad \theta \in [0, 1]
 \end{aligned}$$

Conclusion: the posterior \sim beta distribution($k+r, n-k+s$).

Recall that the prior \sim beta distribution(r, s).

It remains to determine the values of r and s to incorporate the prior knowledge:

PK 1. Mean is about 0.035.

$$\mathbb{E}(\Theta) = 0.035 \quad \implies \quad \frac{r}{r+s} = 0.035 \quad \iff \quad \frac{r}{s} = \frac{7}{193}$$


PK 2. The pdf mostly concentrated over $[0, 0.07]$ trial ...

```

1 x <- seq(0, 1, length = 1025)
2 plot(x, dbeta(x, 4, 102),
3      type="l")
4 plot(x, dbeta(x, 7, 193),
5      type="l")
6 dev.off()

7
8 pdf=cbind(dbeta(x, 4, 102), dbeta(x, 7, 193))
9 matplot(x, pdf,
10         type="l",
11         lty = 1:2,
12         xlab = "theta", ylab = "PDF",
13         lwd = 2 # Line width
14         )
15 legend(0.2, 25, # Position of legend
16        c("Beta(4,102)", "Beta(7,193)"),
17        col = 1:2, lty = 1:2,
18        ncol = 1, # Number of columns
19        cex = 1.5, # Fontsize
20        lwd=2 # Line width
21        )
22 abline(v=0.07, col="blue", lty=1, lwd=1.5)
23 text(0.07, -0.5, "0.07")
24 abline(v=0.035, col="gray60", lty=3, lwd=2)
25 text(0.035, 1, "0.035")

```



If we choose $r = 7$ and $s = 193$:

$$g_{\Theta}(\theta|X = k) = \frac{\Gamma(n + 200)}{\Gamma(k + 7)\Gamma(n - k + 193)} \theta^{k+6} (1 - \theta)^{n-k+192}, \quad \theta \in [0, 1]$$


Moreover, if $n = 10$ and $k = 2$,

$$g_{\Theta}(\theta|X = k) = \frac{\Gamma(210)}{\Gamma(9)\Gamma(201)} \theta^8 (1 - \theta)^{200}, \quad \theta \in [0, 1]$$

```

1 x <- seq(0, 0.1, length = 1025)
2 pdf=cbind(dbeta(x,7,193),dbeta(x,9,201))
3 matplot(x,pdf,
4         type="l",
5         lty = 1:2,
6         xlab = "theta", ylab = "PDF",
7         lwd = 2 # Line width
8     )
9 legend(0.05, 25, # Position of legend
10      c("Posterior Beta(9,201)", "Prior
11        Beta(7,193)"),
12      col = 1:2, lty = 1:2,
13      ncol = 1, # Number of columns
14      cex = 1.5, # Fontsize
15      lwd=2 # Line width
16  )
17 abline(v=0.07,col="blue", lty=1,lwd=1.5)
18 text(0.07, -0.5, "0.07")
19 abline(v=0.035,col="black", lty=3,lwd=2)
20 text(0.035, 1, "0.035")

```



Definition. If the posterior distributions $p(\Theta|X)$ are in the same probability distribution family as the prior probability distribution $p(\Theta)$, the prior and posterior are then called **conjugate distributions**, and the prior is called a **conjugate prior** for the likelihood function.

1. Beta distributions are conjugate priors for Bernoulli, binomial, nega. binomial, geometric likelihood.
2. Gamma distributions are conjugate priors for Poisson and exponential likelihood.

E.g. 2. Let X_1, \dots, X_n be a random sample from $\text{Poisson}(\theta)$: $p_X(k; \theta) = \frac{e^{-\theta} \theta^k}{k!}$ for $k = 0, 1, \dots$.

Let $W = \sum_{i=1}^n X_i$. Then W follows $\text{Poisson}(n\theta)$.

Prior distribution: $\Theta \sim \text{Gamma}(s, \mu)$, i.e., $f_\Theta(\theta) = \frac{\mu^s}{\Gamma(s)} \theta^{s-1} e^{-\mu\theta}$ for $\theta > 0$.

$$\begin{aligned} X_1, \dots, X_n \mid \theta &\sim \text{Poisson}(\theta) \\ \Theta &\sim \text{Gamma}(s, \mu) \\ s \text{ \& } \mu &\text{ are known} \end{aligned}$$

$$\begin{aligned} W = \sum_{i=1}^n X_i \mid \theta &\sim \text{Poisson}(n\theta) \\ \Theta &\sim \text{Gamma}(s, \mu) \\ s \text{ \& } \mu &\text{ are known} \end{aligned}$$

$$g_{\Theta}(\theta|W = w) = \frac{p_W(w|\Theta = \theta)f_{\Theta}(\theta)}{\int_{\mathbb{R}} p_W(w|\Theta = \theta')f_{\Theta}(\theta')d\theta'}$$

$$\begin{aligned} p_W(w|\Theta = \theta)f_{\Theta}(\theta) &= \frac{e^{-n\theta}(n\theta)^w}{w!} \times \frac{\mu^s}{\Gamma(s)}\theta^{s-1}e^{-\mu\theta} \\ &= \frac{n^w}{w!} \frac{\mu^s}{\Gamma(s)} \times \theta^{w+s-1}e^{-(\mu+n)\theta}, \quad \theta > 0. \end{aligned}$$

$$\begin{aligned} p_W(w) &= \int_{\mathbb{R}} p_W(w|\Theta = \theta')f_{\Theta}(\theta')d\theta' \\ &= \frac{n^w}{w!} \frac{\mu^s}{\Gamma(s)} \int_0^{\infty} \theta'^{w+s-1}e^{-(\mu+n)\theta'} d\theta' \\ &= \frac{n^w}{w!} \frac{\mu^s}{\Gamma(s)} \times \frac{\Gamma(w+s)}{(\mu+n)^{w+s}} \end{aligned}$$

$$\begin{aligned}
 g_{\Theta}(\theta|X = k) &= \frac{\frac{n^w}{w!} \frac{\mu^s}{\Gamma(s)} \times \theta^{w+s-1} e^{-(\mu+n)\theta}}{\frac{n^w}{w!} \frac{\mu^s}{\Gamma(s)} \times \frac{\Gamma(w+s)}{(\mu+n)^{w+s}}} \\
 &= \frac{(\mu+n)^{w+s}}{\Gamma(w+s)} \theta^{w+s-1} e^{-(\mu+n)\theta}, \quad \theta > 0.
 \end{aligned}$$

Conclusion: the posterior of $\Theta \sim \text{gamma distribution}(w + s, n + \mu)$.

Recall that the prior of $\Theta \sim \text{gamma distribution}(s, \mu)$.

Case Study 5.8.1


```

1 x <- seq(0, 4, length = 1025)
2 pdf=cbind(dgamma(x, shape=88, rate=50),
3           dgamma(x, shape=88+92, 100),
4           dgamma(x, 88+92+72, 150))
5 matplot(x,pdf,
6         type="l",
7         lty = 1:3,
8         xlab = "theta", ylab = "PDF",
9         lwd = 2 # Line width
10        )
11 legend(2, 3.5, # Position of legend
12       c("Prior Gamma(88,50)",
13         "Posterior1 Beta(180,100)",
14         "Posterior2 Beta(252,150)"),
15       col = 1:3, lty = 1:3,
16       ncol = 1, # Number of columns
17       cex = 1.5, # Fontsize
18       lwd=2 # Line width
19      )

```

Table 5.8.1

Years	Number of Hurricanes
1851–1900	88
1901–1950	92
1951–2000	72



Bayesian Point Estimation

Question. Can one calculate an appropriate *point estimate* θ_e given the posterior $g_{\Theta}(\theta|W = w)$?

Definitions. Let θ_e be an estimate for θ based on a statistic W . The **loss function** associated with θ_e is denoted $L(\theta_e, \theta)$, where $L(\theta_e, \theta) \geq 0$ and $L(\theta, \theta) = 0$.

Let $g_{\Theta}(\theta|W = w)$ be the posterior distribution of the random variable Θ . Then the **risk** associated with $\hat{\theta}$ is **the expected value of the loss function** with respect to the posterior distribution of Θ :


$$\text{risk} = \begin{cases} \int_{\mathbb{R}} L(\hat{\theta}, \theta) g_{\Theta}(\theta|W = w) d\theta & \text{if } \Theta \text{ is continuous} \\ \sum_i L(\hat{\theta}, \theta_i) g_{\Theta}(\theta_i|W = w) & \text{if } \Theta \text{ is discrete} \end{cases}$$

Theorem. Let $g_{\Theta}(\theta|W = w)$ be the posterior distribution of the random variable Θ .


1. If $L(\theta_e, \theta) = |\theta_e - \theta|$, then the Bayes point estimate for θ is the **median** of $g_{\Theta}(\theta|W = w)$.
2. If $L(\theta_e, \theta) = (\theta_e - \theta)^2$, then the Bayes point estimate for θ is the **mean** of $g_{\Theta}(\theta|W = w)$.

Remarks


1. **Median** usually does not have a closed form formula.
2. **Mean** usually has a closed formula.



mode



median



mean

<https://en.wikipedia.org>

Proof. (of Part 1.)

Let m be the *median* of the random variable W . We first claim that

$$\mathbb{E}(|W - m|) \leq \mathbb{E}(|W|). \quad (\star)$$

For any constant $b \in \mathbb{R}$, because

$$\frac{1}{2} = \mathbb{P}(W \leq m) = \mathbb{P}(W - b \leq m - b)$$

we see that $m - b$ is the *median* of $W - b$. Hence, by (\star) ,

$$\mathbb{E}(|W - m|) = \mathbb{E}(|(W - b) - (m - b)|) \leq \mathbb{E}(|W - b|), \quad \text{for all } b \in \mathbb{R},$$

which proves the statement.

Proof. (of Part 1. continued)

It remains to prove (\star) . Without loss of generality, we may assume $m > 0$. Then

$$\begin{aligned}\mathbb{E}(|W - m|) &= \int_{\mathbb{R}} |w - m| f_W(w) dw \\&= \int_{-\infty}^m (m - w) f_W(w) dw + \int_m^{\infty} (w - m) f_W(w) dw \\&= - \int_{-\infty}^m w f_W(w) dw + \int_m^{\infty} w f_W(w) dw + \frac{1}{2}(m - m) \\&= - \int_{-\infty}^0 w f_W(w) dw - \underbrace{\int_0^m w f_W(w) dw}_{\geq 0} + \int_m^{\infty} w f_W(w) dw \\&\leq - \int_{-\infty}^0 w f_W(w) dw + \int_0^{\infty} w f_W(w) dw \\&= \int_{\mathbb{R}} |w| f_W(w) dw \\&= \mathbb{E}(|W|).\end{aligned}$$

Proof. (of Part 2.)

Let μ be the **mean** of W . Then for any $b \in \mathbb{R}$, we see that

$$\begin{aligned}\mathbb{E} [(W - b)^2] &= \mathbb{E} [(W - \mu) + (\mu - b)]^2 \\&= \mathbb{E} [(W - \mu)^2] + 2(\mu - b) \underbrace{\mathbb{E}(W - \mu)}_{=0} + (\mu - b)^2 \\&= \mathbb{E} [(W - \mu)^2] + (\mu - b)^2 \\&\geq \mathbb{E} [(W - \mu)^2] ,\end{aligned}$$

that is,

$$\mathbb{E} [(W - \mu)^2] \leq \mathbb{E} [(W - b)^2] , \quad \text{for all } b \in \mathbb{R}.$$



E.g. 1'. $X_1, \dots, X_n | \theta \sim \text{Bernoulli}(\theta)$ $X = \sum_{i=1}^n X_i | \theta \sim \text{Binomial}(n, \theta)$

$\Theta \sim \text{Beta}(r, s)$ $\Theta \sim \text{Beta}(r, s)$


r & s are known r & s are known

Prior $\text{Beta}(r, s) \rightarrow$ posterior $\text{Beta}(k + r, n - k + s)$
upon observing $X = k$ for a random sample of size n .

Consider the L^2 loss function.

$$\begin{aligned} \theta_e &= \text{mean of } \text{Beta}(k + r, n - k + s) \\ &= \frac{k + r}{n + r + s} \\ &= \frac{n}{n + r + s} \times \underbrace{\left(\frac{k}{n}\right)}_{\text{MLE}} + \frac{r + s}{n + r + s} \times \underbrace{\left(\frac{r}{r + s}\right)}_{\text{Mean of Prior}} \end{aligned}$$

MLE vs. Prior



θ_e

||

$$\underbrace{\frac{n}{n+r+s} \times \left(\frac{k}{n}\right)}_{\text{MLE}} + \underbrace{\frac{r+s}{n+r+s} \times \left(\frac{r}{r+s}\right)}_{\text{Mean of Prior}}$$


E.g. 2'. $X_1, \dots, X_n \mid \theta \sim \text{Poisson}(\theta)$ $W = \sum_{i=1}^n X_i \mid \theta \sim \text{Poisson}(n\theta)$
 $\Theta \sim \text{Gamma}(s, \mu)$ $\Theta \sim \text{Gamma}(s, \mu)$
 $s \text{ \& } \mu \text{ are known}$ $s \text{ \& } \mu \text{ are known}$

Prior $\text{Gamma}(s, \mu) \rightarrow$ Posterior $\text{Gamma}(w + s, \mu + n)$
upon observing $W = w$ for a random sample of size n .

Consider the L^2 loss function.

$$\begin{aligned} \theta_e &= \text{mean of Gamma}(w + s, \mu + n) \\ &= \frac{w + s}{\mu + n} \\ &= \frac{n}{\mu + n} \times \underbrace{\left(\frac{w}{n}\right)}_{\text{MLE}} + \frac{\mu}{\mu + n} \times \underbrace{\left(\frac{s}{\mu}\right)}_{\text{Mean of Prior}} \end{aligned}$$

MLE vs. Prior



θ_e

||

$$\frac{n}{\mu + n} \times \underbrace{\left(\frac{w}{n}\right)}_{\text{MLE}} + \frac{\mu}{\mu + n} \times \underbrace{\left(\frac{s}{\mu}\right)}_{\text{Mean of Prior}}$$

Appendix: Beta integral

Lemma. $B(\alpha, \beta) := \int_0^1 x^{\alpha-1} (1-x)^{\beta-1} dx = \frac{\Gamma(\alpha)\Gamma(\beta)}{\Gamma(\alpha+\beta)}$

Proof. Notice that

$$\Gamma(\alpha) = \int_0^\infty x^{\alpha-1} e^{-x} dx \quad \text{and} \quad \Gamma(\beta) = \int_0^\infty y^{\beta-1} e^{-y} dy.$$

Hence,

$$\Gamma(\alpha)\Gamma(\beta) = \int_0^\infty \int_0^\infty x^{\alpha-1} y^{\beta-1} e^{-(x+y)} dx dy.$$

The key in the proof is the following change of variables:

$$\begin{cases} x = r^2 \cos^2(\theta) \\ y = r^2 \sin^2(\theta) \end{cases}$$

$$\implies \frac{\partial(x, y)}{\partial(r, \theta)} = \begin{pmatrix} 2r \cos^2(\theta) & 2r \sin^2(\theta) \\ -2r^2 \cos(\theta) \sin(\theta) & 2r^2 \cos(\theta) \sin(\theta) \end{pmatrix}$$

$$\implies \left| \det \left(\frac{\partial(x, y)}{\partial(r, \theta)} \right) \right| = 4r^3 \sin(\theta) \cos(\theta).$$

Therefore,

$$\begin{aligned}\Gamma(\alpha)\Gamma(\beta) &= \int_0^{\frac{\pi}{2}} d\theta \int_0^\infty dr r^{2(\alpha+\beta)-4} e^{-r^2} \cos^{2\alpha-2}(\theta) \sin^{2\beta-2}(\theta) \times \underbrace{4r^3 \sin(\theta) \cos(\theta)}_{\text{Jacobian}} \\ &= 4 \left(\int_0^{\frac{\pi}{2}} \cos^{2\alpha-1}(\theta) \sin^{2\beta-1}(\theta) d\theta \right) \left(\int_0^\infty r^{2(\alpha+\beta)-1} e^{-r^2} dr \right).\end{aligned}$$

Now let us compute the following two integrals separately:

$$I_1 := \int_0^{\frac{\pi}{2}} \cos^{2\alpha-1}(\theta) \sin^{2\beta-1}(\theta) d\theta$$

$$I_2 := \int_0^\infty r^{2(\alpha+\beta)-1} e^{-r^2} dr$$

For I_2 , by change of variable $r^2 = u$ (so that $2rdr = du$),

$$\begin{aligned} I_2 &= \int_0^\infty r^{2(\alpha+\beta)-1} e^{-r^2} dr \\ &= \frac{1}{2} \int_0^\infty r^{2(\alpha+\beta)-2} e^{-r^2} \underbrace{2rdr}_{=du} \\ &= \frac{1}{2} \int_0^\infty u^{\alpha+\beta-1} e^{-u} du \\ &= \frac{1}{2} \Gamma(\alpha + \beta). \end{aligned}$$

For I_1 , by the change of variables $\sqrt{x} = \cos(\theta)$ (so that $-\sin(\theta)d\theta = \frac{1}{2\sqrt{x}}dx$),

$$\begin{aligned}
 I_1 &= \int_0^{\frac{\pi}{2}} \cos^{2\alpha-1}(\theta) \sin^{2\beta-1}(\theta) d\theta \\
 &= \int_0^{\frac{\pi}{2}} \cos^{2\alpha-1}(\theta) \sin^{2\beta-2}(\theta) \times \underbrace{\sin(\theta)d\theta}_{=-\frac{1}{2\sqrt{x}}dx} \\
 &= \int_1^0 x^{\alpha-\frac{1}{2}} (1-x)^{\beta-1} \frac{-1}{2\sqrt{x}} dx \\
 &= \frac{1}{2} \int_0^1 x^{\alpha-1} (1-x)^{\beta-1} dx \\
 &= \frac{1}{2} B(\alpha, \beta)
 \end{aligned}$$

Therefore,

$$\begin{aligned}\Gamma(\alpha)\Gamma(\beta) &= 4l_1 \times l_2 \\ &= 4 \times \frac{1}{2}\Gamma(\alpha + \beta) \times \frac{1}{2}B(\alpha, \beta)\end{aligned}$$

i.e.,

$$B(\alpha, \beta) = \frac{\Gamma(\alpha)\Gamma(\beta)}{\Gamma(\alpha + \beta)}.$$

□